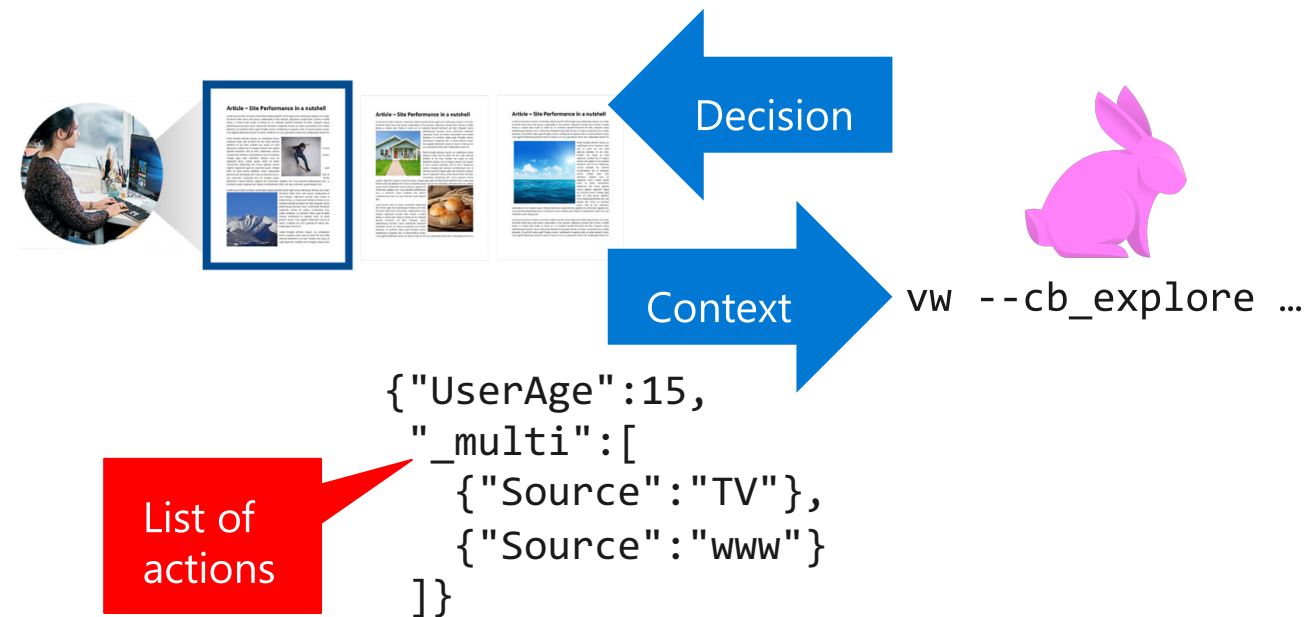


VW Workshop Slates

Adith Swaminathan & Jack Gerrits
Microsoft Research

cb_explore: Contextual Decisions with Exploration

- Use contextual bandit algorithms to pick actions that maximize rewards
 - Balances exploration to discover rewarding actions against exploitation
 - Tunable parameters: exploration rate, learning rate, base learner type, etc.
- Action set represented as a list
 - Either a fixed set across contexts, or,
 - Featurized action lists per context



Slates: Motivating Applications



Network configuration



Page layout optimization



Slate recommendations

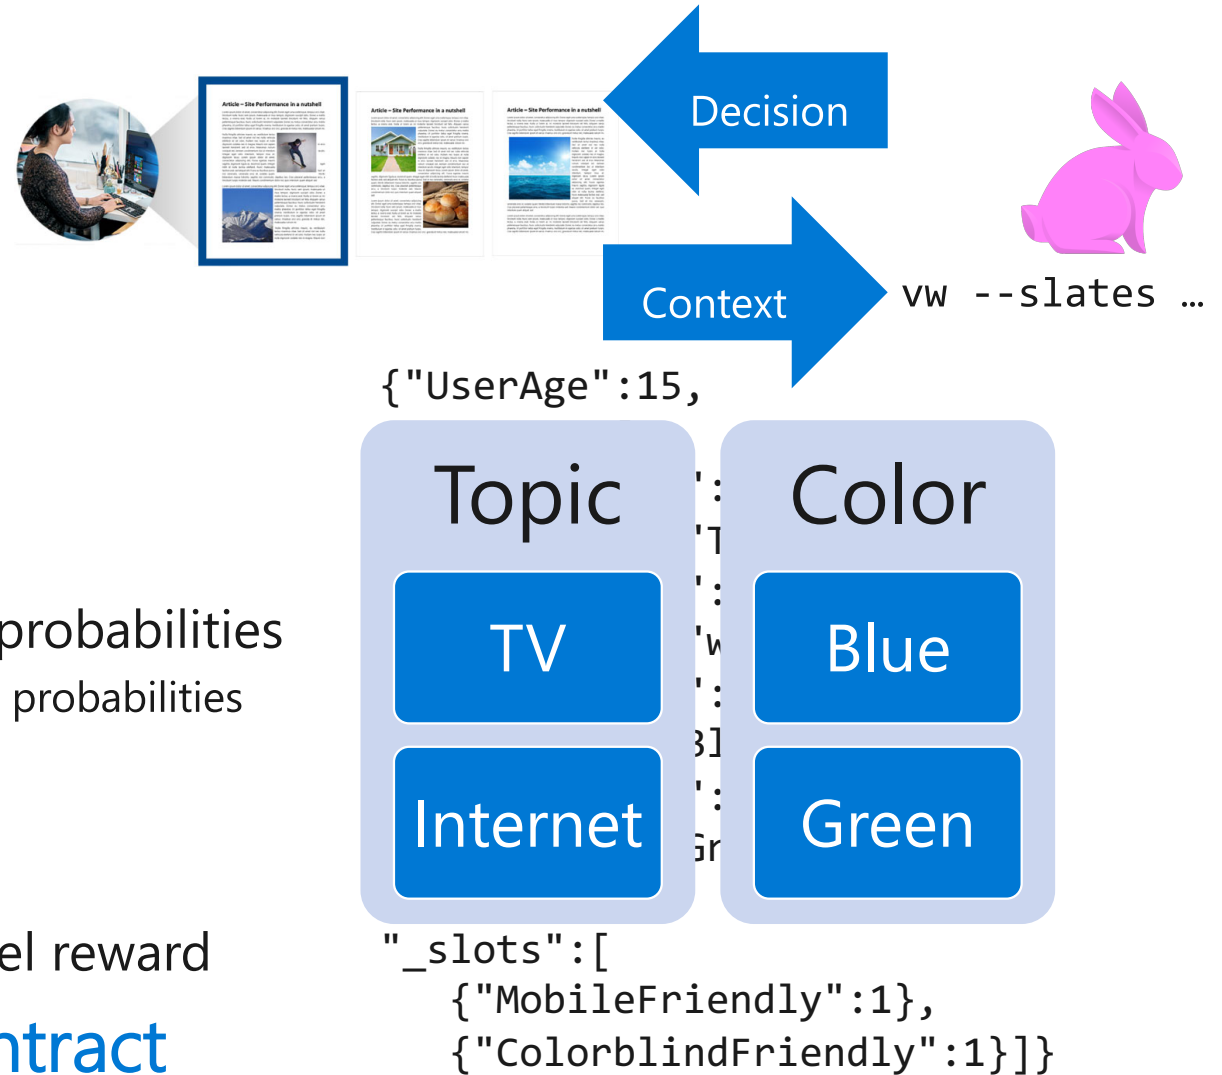
- “Composite” actions that are composed of several choices
 - Each choice populates a “slot”
 - Reward is observed for the composite action
 - **Combinatorially many** actions if encoded as a list
- VW automagically handles feature interactions across namespaces...
- Can it handle these action-slot-interactions also? **Yes! --slates**

Slates: Problem Definition

- Context
 - Shared features + Number of slots
 - Slot-specific features
 - Action-set for each slot
 - **Disjoint** action-sets across slots
- Decision
 - List (len=#slots) of action ids and selection probabilities
 - **Semantics**: probability of slate = product of selection probabilities
 - Action ids index into respective action set
- Reward
 - A single scalar representing global slate-level reward

Contract

VW will run cb explore-exploit with actions as the cartesian product of slot-specific action sets



Under The Hood

VW --slates repeatedly calls `cb_explore` with book-keeping

- One `cb_explore` instance for each slot
- Context automatically computes correct shared/slot/action feature interactions
- Action constructed by invoking each `cb_explore` instance to populate slots independently
- Reward at slate-level is automatically decomposed into slot-specific rewards for each learner

Tunable hyper-parameters: Same as for base learner (`cb_explore`).

- No additional slate-specific tuning parameters

`slates` advantages over `cb_explore`:

1. Efficient encoding
2. Faster learning
3. Correct feature interactions

When to use slates vs. cb_explore vs. ccb?

cb_explore: General purpose base-learner. Works for small #actions.

ccb: Extends cb_explore to slate recommendation scenarios.

- Expects slot-specific rewards. (No slate->slot credit assignment)
- Works with dependent slot-specific action sets. (Supports rankings)

slates: Extends cb_explore to slate recommendation scenarios.

- Works with only slate-level reward. (Performs slate->slot credit assignment)
- Currently supports only disjoint slot-specific action sets. (E.g., no rankings support yet)

<https://arxiv.org/abs/1605.04812>

A Concrete Example & Demo

Outfit Optimization for Job Interview



https://www.youtube.com/watch?v=G-omu_ki7YM